

19075153_O'Leary_PartB2

September 11, 2021

```
[1]: # Feature Extraction with Univariate Statistical Tests (Chi-squared for
      ↪ classification)
import pandas
import numpy as np
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.feature_selection import RFE
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score
from sklearn.decomposition import PCA
```

1 Data pre-processing performed

```
[2]: # load data
path = "Autism-Screening-Child-Data Plus Description/Autism-Child-Data.arff"
from scipy.io.arff import loadarff
raw_data = loadarff(path)
df_data = pandas.DataFrame(raw_data[0])

# Remove the b'' from the data
str_df = df_data.select_dtypes([np.object])
str_df = str_df.stack().str.decode('utf-8').unstack()
for col in str_df:
    df_data[col] = str_df[col]

# Make the whole lot categorical
list_of_columns = df_data.columns
df_data[list_of_columns] = df_data[list_of_columns].apply(lambda col:pandas.
    ↪Categorical(col).codes)

# Remove all rows with negative values
df_data = df_data[(df_data >= 0).all(1)]
```

```
# split the data to X and y
X = df_data.loc[:, df_data.columns != 'result']
Y = df_data.iloc[:, 17]
#with pandas.option_context('display.max_rows', None, 'display.max_columns',
↳None):
#    print(df_data)
df_data
```

<ipython-input-2-b56ce8af1f47>:8: DeprecationWarning: `np.object` is a deprecated alias for the builtin `object`. To silence this warning, use `object` by itself. Doing this will not modify any behavior and is safe.
Deprecated in NumPy 1.20; for more details and guidance:
<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>
str_df = df_data.select_dtypes([np.object])

```
[2]:
```

	A1_Score	A2_Score	A3_Score	A4_Score	A5_Score	A6_Score	A7_Score	\
0	1	1	0	0	1	1	0	
1	1	1	0	0	1	1	0	
2	1	1	0	0	0	1	1	
3	0	1	0	0	1	1	0	
4	1	1	1	1	1	1	1	
..	
287	1	1	1	1	1	1	1	
288	1	0	0	0	1	0	1	
289	1	0	1	1	1	1	1	
290	1	1	1	0	1	1	1	
291	0	0	1	0	1	0	1	

	A8_Score	A9_Score	A10_Score	...	gender	ethnicity	jundice	austim	\
0	1	0	0	...	1	6	0	0	
1	1	0	0	...	1	5	0	0	
2	1	0	0	...	1	0	0	0	
3	0	0	1	...	0	0	1	0	
4	1	1	1	...	1	6	1	0	
..	
287	1	1	1	...	0	10	1	1	
288	0	0	1	...	0	10	1	1	
289	0	0	1	...	1	4	0	0	
290	1	1	1	...	1	8	0	0	
291	0	0	0	...	0	8	0	0	

	contry_of_res	used_app_before	result	age_desc	relation	Class/ASD
0	24	0	5	0	2	0
1	24	0	5	0	2	0
2	24	1	5	0	0	0
3	24	0	4	0	0	0
4	51	0	10	0	2	1

..
287	50	0	10	0	2	1
288	3	0	4	0	2	0
289	8	0	7	0	2	1
290	18	0	9	0	2	1
291	18	0	3	0	2	0

[288 rows x 21 columns]

2 Feature selection

```
[3]: # generate model and get accuracy
def get_accuracy(target_train, target_test, predicted_test, predicted_train):
    clf = MLPClassifier(activation='logistic', solver='sgd',
        ↪learning_rate_init=0.1, alpha=1e-5, hidden_layer_sizes=(5, 2),
        ↪random_state=1,max_iter=2000)
    clf.fit(predicted_train, np.ravel(target_train, order='C'))
    predictions = clf.predict(predicted_test)
    return accuracy_score(target_test, predictions)

pred_train, pred_test, tar_train, tar_test = train_test_split(X, Y, test_size=.
    ↪3, random_state=4)
print("Accuracy score of our model without feature selection : %.2f" %
    ↪get_accuracy(tar_train, tar_test, pred_test, pred_train))

# feature extraction
test = SelectKBest(score_func=chi2, k=5)
fit = test.fit(X, Y)
# summarize scores
np.set_printoptions(precision=3)
print(fit.scores_)
features = fit.transform(X)
# summarize selected features
print(features[0:5, :], "summerize features")
print()
# Now apply only the K most significant features according to the chi square
    ↪method
pred_features = features[:, 0:5]
pred_train, pred_test, tar_train, tar_test = train_test_split(pred_features, Y,
    ↪test_size=.3, random_state=2)
print("Accuracy score of our model with chi square feature selection : %.2f" %
    ↪get_accuracy(tar_train, tar_test, pred_test, pred_train))
print()

## Feature Importance with Extra Trees Classifier
from sklearn.feature_selection import SelectFromModel
```

```

# Feature Extraction with RFE
model = LogisticRegression() # Logistic regression is the Wrapper classifier
    ↳ here
rfe = RFE(model, 5)
fit = rfe.fit(X, Y)
## summarize components
#print("Num Features: %d" % (fit.n_features_))
#print("Selected Features: %s" % (fit.support_))
#print("Feature Ranking: %s" % (fit.ranking_))
## Now apply only the K most significant features according to the RFE feature
    ↳ selection method
features = fit.transform(X)
pred_features = features[:, 0:5]
pred_train, pred_test, tar_train, tar_test = train_test_split(pred_features, Y,
    ↳ test_size=.3, random_state=2)
print("Accuracy score of our model with RFE selection : %.2f" %
    ↳ get_accuracy(tar_train, tar_test, pred_test, pred_train))
print()

## Feature Extraction with PCA
## feature extraction
pca = PCA(n_components=5)
fit = pca.fit(X)
features = fit.transform(X)
## summarize components
#print("Explained Variance: %s" % (fit.explained_variance_ratio_))
#print(fit.components_)
## Now apply only the K most significant faetures (components) according to the
    ↳ PCA feature selection method
#features = fit.transform(X)
pred_features = features[:, 0:5]
pred_train, pred_test, tar_train, tar_test = train_test_split(pred_features, Y,
    ↳ test_size=.3, random_state=2)
print("Accuracy score of our model with PCA selection : %.2f" %
    ↳ get_accuracy(tar_train, tar_test, pred_test, pred_train))
print()

## Feature Importance with Extra Trees Classifier
from sklearn.ensemble import ExtraTreesClassifier
## feature extraction
model = ExtraTreesClassifier(max_depth=3, min_samples_leaf=2)
fit = model.fit(X, Y)
print(model.feature_importances_)
print()
t = SelectFromModel(fit, prefit=True)
features = t.transform(X)

```

```

pred_features = features[:, 0:5]
pred_train, pred_test, tar_train, tar_test = train_test_split(pred_features, Y,
    ↳test_size=.3, random_state=2)
print("Accuracy score of our model with Extra Trees selection : %.2f" %
    ↳get_accuracy(tar_train, tar_test, pred_test, pred_train))
print()

```

```

Accuracy score of our model without feature selection : 0.05
[ 20.321  22.921  21.164  51.545  18.469  33.979  18.496  43.588  56.682
  20.841  26.617   1.917  56.078   5.255   5.473 121.11   1.619    nan
   4.792 149.    ]
[[ 0  0  6 24  0]
 [ 0  0  5 24  0]
 [ 0  0  0 24  0]
 [ 0  0  0 24  0]
 [ 1  1  6 51  1]] summerize features

```

Accuracy score of our model with chi square feature selection : 0.15

```

/home/bernard/anaconda3/lib/python3.8/site-
packages/sklearn/utils/validation.py:70: FutureWarning: Pass
n_features_to_select=5 as keyword args. From version 1.0 (renaming of 0.25)
passing these as positional arguments will result in an error
  warnings.warn(f"Pass {args_msg} as keyword args. From version "
/home/bernard/anaconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:763: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```

n_iter_i = _check_optimize_result(
/home/bernard/anaconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:763: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```

n_iter_i = _check_optimize_result(
/home/bernard/anaconda3/lib/python3.8/site-

```

```
packages/sklearn/linear_model/_logistic.py:763: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/home/bernard/anaconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:763: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/home/bernard/anaconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:763: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/home/bernard/anaconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:763: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/home/bernard/anaconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:763: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/home/bernard/anaconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:763: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/home/bernard/anaconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:763: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/home/bernard/anaconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:763: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
/home/bernard/anaconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:763: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```

regression
    n_iter_i = _check_optimize_result(
/home/bernard/anaconda3/lib/python3.8/site-
packages/sklearn/linear_model/_logistic.py:763: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```

regression
    n_iter_i = _check_optimize_result(

```

Accuracy score of our model with RFE selection : 0.15

Accuracy score of our model with PCA selection : 0.15

```

[0.026 0.036 0.053 0.115 0.043 0.091 0.04  0.059 0.133 0.038 0.008 0.003
 0.007 0.002 0.002 0.003 0.001 0.    0.01  0.329]

```

Accuracy score of our model with Extra Trees selection : 0.15

3 Get five most important features

```

[4]: from sklearn.ensemble import ExtraTreesClassifier
import matplotlib.pyplot as plt

# Get the average of the list
def Average(lst):
    return sum(lst) / len(lst)


## feature extraction
model = ExtraTreesClassifier(max_depth=3,min_samples_leaf=2)
fit = model.fit(X, Y)
print(model.feature_importances_)
print()
feat_importances = pandas.Series(model.feature_importances_, index=X.columns)
feat_importances.nlargest(20).plot(kind='barh')

t = SelectFromModel(fit, prefit=True)
features = t.transform(X)
pred_features = features[:, 0:5]
pred_train, pred_test, tar_train, tar_test = train_test_split(pred_features, Y,
↳test_size=.3, random_state=4)

```



```
[0.017 0.038 0.038 0.114 0.036 0.107 0.043 0.049 0.136 0.035 0.006 0.004
 0.012 0.002 0.001 0.002 0.    0.    0.008 0.355]
```



4 Naive Bayes Model

```
[5]: from sklearn.model_selection import KFold
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import precision_score, recall_score, auc, roc_curve

from sklearn import svm, datasets
from sklearn.model_selection import train_test_split
from sklearn.metrics import plot_confusion_matrix

gnb = GaussianNB() #suitable for numeric features
gnb.fit(pred_train, np.ravel(tar_train, order='C'))
kf = KFold(n_splits=10)
gnb.fit(pred_train, tar_train)

# Plot non-normalized confusion matrix
titles_options = [("Confusion matrix, without normalization", None),
                  ("Normalized confusion matrix", 'true')]
for title, normalize in titles_options:
    disp = plot_confusion_matrix(gnb, pred_test, tar_test,
                                #display_labels=['yes', 'no'],
```

```

                                cmap=plt.cm.Blues,
                                normalize=normalize)

disp.ax_.set_title(title)

print(title)
print(disp.confusion_matrix)

plt.xlabel('Predictions', fontsize=18)
plt.ylabel('Actuals', fontsize=18)
plt.show()

# Print accuracy
print(Average(cross_val_score(gnb, pred_train, tar_train, cv=kf,
→scoring='accuracy'))))

```

Confusion matrix, without normalization

```

[[ 0  1  0  0  0  0  0  0  0  0  0]
 [ 0  0  0  0  0  0  0  0  0  0  0]
 [ 0  2  1  0  0  1  1  0  0  0  0]
 [ 0  3  0  0  0  0  1  0  0  0  0]
 [ 0  4  0  0  0  2  1  0  0  0  0]
 [ 0  3  1  0  0  6  4  0  0  0  0]
 [ 0  0  0  0  0  8  2  0  0  0  0]
 [ 0  0  0  0  0  0  0 14  5  0  3]
 [ 0  0  0  0  0  0  0  5  2  0  7]
 [ 0  0  0  0  0  0  0  0  0  0  4]
 [ 0  0  0  0  0  0  0  0  0  0  6]]

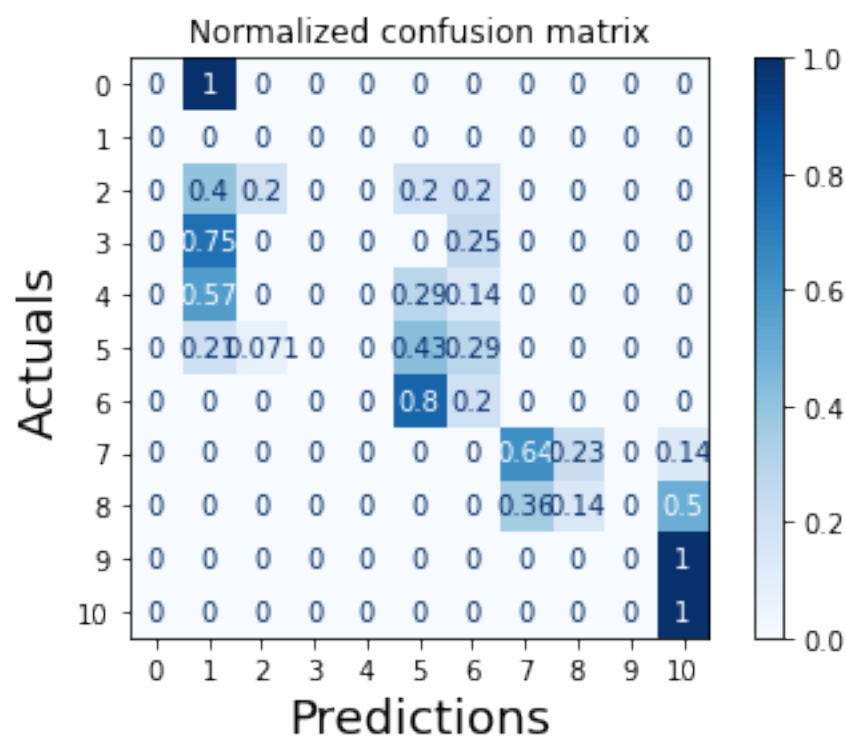

```

Normalized confusion matrix

```

[[0.    1.    0.    0.    0.    0.    0.    0.    0.    0.    0. ]
 [0.    0.    0.    0.    0.    0.    0.    0.    0.    0.    0. ]
 [0.    0.4  0.2  0.    0.    0.2  0.2  0.    0.    0.    0. ]
 [0.    0.75 0.    0.    0.    0.    0.25 0.    0.    0.    0. ]
 [0.    0.571 0.    0.    0.    0.286 0.143 0.    0.    0.    0. ]
 [0.    0.214 0.071 0.    0.    0.429 0.286 0.    0.    0.    0. ]
 [0.    0.    0.    0.    0.    0.8  0.2  0.    0.    0.    0. ]
 [0.    0.    0.    0.    0.    0.    0.    0.636 0.227 0.    0.136]
 [0.    0.    0.    0.    0.    0.    0.    0.357 0.143 0.    0.5  ]
 [0.    0.    0.    0.    0.    0.    0.    0.    0.    0.    1.  ]
 [0.    0.    0.    0.    0.    0.    0.    0.    0.    0.    1.  ]]

```



0.35261904761904767

5 Decision Tree Classifier Model

```
[6]: from sklearn.tree import DecisionTreeClassifier
import numpy as np
import matplotlib.pyplot as plt

from sklearn import svm, datasets
from sklearn.model_selection import train_test_split
from sklearn.metrics import plot_confusion_matrix

# Reset the train/test-split
pred_train, pred_test, tar_train, tar_test = train_test_split(X, Y, test_size=.
    ↪3, random_state=4)

# Train the DTC
dt = DecisionTreeClassifier(min_samples_leaf=0.2, min_samples_split=0.3)
kf = KFold(n_splits=10)
dt.fit(pred_train, tar_train)
train_pred = dt.predict(pred_train)

# Plot non-normalized confusion matrix
titles_options = [("Confusion matrix, without normalization", None),
    ("Normalized confusion matrix", 'true')]
for title, normalize in titles_options:
    disp = plot_confusion_matrix(dt, pred_test, tar_test,
        #display_labels=['yes', 'no'],
        cmap=plt.cm.Blues,
        normalize=normalize)

    disp.ax_.set_title(title)

    print(title)
    print(disp.confusion_matrix)

plt.xlabel('Predictions', fontsize=18)
plt.ylabel('Actuals', fontsize=18)
plt.show()

# print accuracy
print(Average(cross_val_score(dt, pred_train, tar_train, cv=kf,
    ↪scoring='accuracy'))))

# get importance
print(dt.feature_importances_)
print()
```


```
feat_importances = pandas.Series(dt.feature_importances_, index=X.columns)
feat_importances.nlargest(20).plot(kind='barh')
```


Confusion matrix, without normalization

```
[[ 0  0  0  1  0  0  0  0  0  0]
 [ 0  0  0  4  0  1  0  0  0  0]
 [ 0  0  0  4  0  0  0  0  0  0]
 [ 0  0  0  4  0  3  0  0  0  0]
 [ 0  0  0  5  0  9  0  0  0  0]
 [ 0  0  0  0  0 10  0  0  0  0]
 [ 0  0  0  0  0  0 15  7  0  0]
 [ 0  0  0  0  0  0  9  5  0  0]
 [ 0  0  0  0  0  0  1  3  0  0]
 [ 0  0  0  0  0  0  3  3  0  0]]
```

Normalized confusion matrix

```
[[0.  0.  0.  1.  0.  0.  0.  0.  0.  0.]
 [0.  0.  0.  0.8 0.  0.2 0.  0.  0.  0.]
 [0.  0.  0.  1.  0.  0.  0.  0.  0.  0.]
 [0.  0.  0.  0.571 0.  0.429 0.  0.  0.  0.]
 [0.  0.  0.  0.357 0.  0.643 0.  0.  0.  0.]
 [0.  0.  0.  0.  0.  1.  0.  0.  0.  0.]
 [0.  0.  0.  0.  0.  0.  0.682 0.318 0.  0.]
 [0.  0.  0.  0.  0.  0.  0.643 0.357 0.  0.]
 [0.  0.  0.  0.  0.  0.  0.25 0.75 0.  0.]
 [0.  0.  0.  0.  0.  0.  0.5  0.5  0.  0.]]
```






0.27809523809523806

```
[0.  0.  0.  0.  0.  0.223 0.  0.  0.  0.  0.  0.
 0.054 0.  0.  0.  0.  0.  0.  0.723]
```

[6]: <AxesSubplot:>



[]: