

19075153_O'Leary_PartA1_DB

October 16, 2021

```
[1]: import pandas
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import warnings
from pandas.plotting import scatter_matrix
import seaborn as sns
from sklearn.model_selection import cross_val_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report, confusion_matrix
from sklearn import tree
from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2
from numpy import set_printoptions
from sklearn.decomposition import PCA
from sklearn import preprocessing

warnings.filterwarnings('ignore')
```

1 Data load and pre-processing

```
[2]: min_max_scaler = preprocessing.MinMaxScaler()

#####
# load data sales
#####
path_sales = "/mnt/c/Users/T828808/Study/AUT/COMP809/Ass2/
↳Sales_Transactions_Dataset_Weekly.csv"
rawdata_sales = pandas.read_csv(path_sales)
# categorise everything and create array
list_of_columns_sales = rawdata_sales.columns
rawdata_sales[list_of_columns_sales] = rawdata_sales[list_of_columns_sales].
↳apply(lambda col:pandas.Categorical(col).codes)
# Create array
array_sales = rawdata_sales.values
predictors_sales = array_sales[:, 0:107]
# Print some stats
```

```

print(rawdata_sales.shape)
print(rawdata_sales.head())

#####
# load data seoul
#####
path_seoul = "/mnt/c/Users/T828808/Study/AUT/COMP809/Ass2/SeoulBikeData.csv"
rawdata_seoul = pandas.read_csv(path_seoul)
# categorise everything and create array
list_of_columns_seoul = rawdata_seoul.columns
rawdata_seoul[list_of_columns_seoul] = rawdata_seoul[list_of_columns_seoul].
    ↳ apply(lambda col:pandas.Categorical(col).codes)
# Create array
array_seoul = rawdata_seoul.values
predictors_seoul = array_seoul[:, 0:14]
# Print some stats
print(rawdata_seoul.shape)
print(rawdata_seoul.head())

#####
# load data water
#####
path_water = "/mnt/c/Users/T828808/Study/AUT/COMP809/Ass2/water-treatment.data"
rawdata_water = pandas.read_csv(path_water)
rawdata_water.columns =
    ↳ ['DATE', 'Q-E', 'ZN-E', 'PH-E', 'DBO-E', 'DQO-E', 'SS-E', 'SSV-E', 'SED-E', 'COND-E', 'PH-P', 'DBO-P',
# categorise everything and create array
list_of_columns_water = rawdata_water.columns
rawdata_water[list_of_columns_water] = rawdata_water[list_of_columns_water].
    ↳ apply(lambda col:pandas.Categorical(col).codes)
# Create array
array_water = rawdata_water.values
predictors_water = array_water[:, 0:39]
# Print some stats
print(rawdata_water.shape)
print(rawdata_water.head())

```

(811, 107)

	Product_Code	W0	W1	W2	W3	W4	W5	W6	W7	W8	...	Normalized 42	\
0	0	11	12	10	8	13	12	14	21	6	...	3	
1	111	7	6	3	2	7	1	6	3	3	...	17	
2	222	7	11	8	9	10	8	7	13	12	...	24	
3	331	12	8	13	5	9	6	9	13	13	...	37	
4	442	8	5	13	11	6	7	9	14	9	...	24	

	Normalized 43	Normalized 44	Normalized 45	Normalized 46	Normalized 47	\
0	20	26	35	46	0	

1	38	47	7	6	35
2	83	16	15	32	40
3	45	4	9	20	30
4	51	25	56	16	15

	Normalized 48	Normalized 49	Normalized 50	Normalized 51
0	16	13	7	35
1	44	7	55	0
2	82	41	41	32
3	65	31	25	31
4	8	49	29	36

[5 rows x 107 columns]

(8760, 14)

	Date	Rented Bike Count	Hour	Temperature(C)	Humidity(%) \
0	11	253	0	111	28
1	11	203	1	108	29
2	11	172	2	103	30
3	11	106	3	101	31
4	11	77	4	103	27

	Wind speed (m/s)	Visibility (10m)	Dew point temperature(C) \
0	22	1788	114
1	8	1788	114
2	10	1788	113
3	9	1788	114
4	23	1788	104

	Solar Radiation (MJ/m2)	Rainfall(mm)	Snowfall (cm)	Seasons	Holiday \
0	0	0	0	3	1
1	0	0	0	3	1
2	0	0	0	3	1
3	0	0	0	3	1
4	0	0	0	3	1

	Functioning Day
0	1
1	1
2	1
3	1
4	1

(526, 39)

	DATE	Q-E	ZN-E	PH-E	DBO-E	DQO-E	SS-E	SSV-E	SED-E	COND-E	...	\
0	197	330	116	6	204	169	57	201	50	409	...	
1	427	99	143	5	204	231	42	208	28	303	...	
2	443	219	126	8	93	256	45	171	36	402	...	
3	461	282	66	9	126	211	37	164	33	381	...	
4	479	319	116	7	90	117	42	197	36	295	...	


	COND-S	RD-DBO-P	RD-SS-P	RD-SED-P	RD-DBO-S	RD-DQO-S	RD-DBO-G	\
0	372	314	165	104	184	233	155	
1	334	314	143	111	184	37	155	
2	322	100	196	108	131	165	95	
3	349	314	183	111	184	153	114	
4	301	314	157	118	125	216	94	

	RD-DQO-G	RD-SS-G	RD-SED-G
0	134	126	0
1	101	92	26
2	158	101	0
3	121	82	37
4	78	61	0


[5 rows x 39 columns]

2 Feature importance


```
[3]: #####
# Sales
#####
# Run PCA
pca_sales = PCA()
pca_fit_sales = pca_sales.fit(predictors_sales)
# summarize components
feat_importances_sales = pandas.Series(pca_fit_sales.explained_variance_ratio_,
→index=rawdata_sales.columns)
feat_importances_sales.nlargest(20).plot(kind='barh')
# Get only the first two components as they explain almost all of the variance
pca_sales = PCA(n_components=2)
pca_fit_sales = pca_sales.fit(predictors_sales)
pca_data_sales = rawdata_sales[['W0', 'W1', 'Product_Code']]
```



```
[4]: #####
# Seoul
#####
# Run PCA
pca_seoul = PCA()
pca_fit_seoul = pca_seoul.fit(predictors_seoul)
# summarize components
feat_importances_seoul = pandas.Series(pca_fit_seoul.explained_variance_ratio_,
    ↪ index=rawdata_seoul.columns)
feat_importances_seoul.nlargest(20).plot(kind='barh')
# Get only the first two components as they explain almost all of the variance
pca_seoul = PCA(n_components=2)
pca_fit_seoul = pca_seoul.fit(predictors_seoul)
pca_data_seoul = rawdata_seoul[['Hour', 'Rented Bike_
    ↪ Count', 'Date', 'Temperature( C)']]
```



```
[5]: #####
# Water
#####
# Run PCA
pca_water = PCA()
pca_fit_water = pca_water.fit(predictors_water)
# summarize components
feat_importances_water = pandas.Series(pca_fit_water.explained_variance_ratio_,
→ index=rawdata_water.columns)
feat_importances_water.nlargest(20).plot(kind='barh')
# Get only the first two components as they explain almost all of the variance
pca_water = PCA(n_components=2)
pca_fit_water = pca_water.fit(predictors_water)
pca_data_water = rawdata_water[['DATE', 'Q-E', 'ZN-E', 'PH-E', 'DBO-E']]
```



3 Clustering

```
[12]: from sklearn.datasets import make_blobs
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_samples, silhouette_score
import matplotlib.pyplot as plt
import matplotlib.cm as cm
import numpy as np
from sklearn.cluster import AgglomerativeClustering
from sklearn.cluster import DBSCAN
import time

def do_sse(X, cluster_labels, n_clusters, model):
    cluster_centers = [X[cluster_labels == i].mean(axis=0) for i in
    ↪range(n_clusters)]
    clusterwise_sse = [0, 0, 0, 0, 0, 0]
    for point, label in zip(X, cluster_labels):
        clusterwise_sse[label] += np.square(point - cluster_centers[label]).
    ↪sum()
    clusterwise_sse_avg = np.mean(clusterwise_sse)
    return clusterwise_sse_avg

def do_cluster_analysis(name):
```

```

    # To find out the optimal number of clusters we can search through range of
    → clusters.
    range_n_clusters = [2, 3, 4, 5, 6]
    for n_clusters in range_n_clusters:

        □
        → print('=====')
            print('n_clusters = ', n_clusters)

        □
        → print('=====')

        start_time = time.time()

        # Create a subplot with 1 row and 2 columns
        fig, (ax1, ax2) = plt.subplots(1, 2)
        fig.set_size_inches(18, 7)

        # The 1st subplot is the silhouette plot
        # The silhouette coefficient can range from -1, 1
        # but in this example code all lie within [-0.1, 1]

        ax1.set_xlim([-0.1, 1])

        # # The (n_clusters+1)*10 is for inserting blank space between
        # silhouette plots of individual clusters, to demarcate them
        # clearly.

        ax1.set_ylim([0, len(X) + (n_clusters + 1) * 10])

        #Apply DBSCAN and set the EPS
        eps = 0
        if name == 'sales':
            eps = 0.2
        elif name == 'water':
            eps = 0.3
        elif name == 'seoul':
            eps = 0.1
        clusterer = DBSCAN(eps, min_samples=n_clusters, metric='euclidean')
        cluster_labels = clusterer.fit_predict(X)

        # The silhouette_score gives the average value for all the
        # samples. This gives a perspective into the density and
        # separation of the formed clusters

        try:
            silhouette_avg = silhouette_score(X, cluster_labels)

```



```

#_
↪ ///////////////////////////////////////////////////
    # Print the values
    #_
↪ ///////////////////////////////////////////////////
    print("For n_clusters =", n_clusters, "The average silhouette_score_
↪is :", silhouette_avg)
    print("For n_clusters =", n_clusters, "The average SSE is :",_
↪do_sse(X, clusterer.labels_, n_clusters, clusterer))
    except:
        print('DBSCAN EXCEPTION')
        break

    # Compute the silhouette scores for each sample
    sample_silhouette_values = silhouette_samples(X, cluster_labels)
    y_lower = 10

    for i in range(n_clusters):
        # Aggregate the silhouette scores for samples belonging to
        # cluster i, and sort them

        #_
↪ ///////////////////////////////////////////////////
        # Create the plot
        #_
↪ ///////////////////////////////////////////////////

        # Aggregate the silhouette scores for samples belonging to
        # cluster i, and sort them
        ith_cluster_silhouette_values =_
↪sample_silhouette_values[cluster_labels == i]

        ith_cluster_silhouette_values.sort()
        size_cluster_i = ith_cluster_silhouette_values.shape[0]
        y_upper = y_lower + size_cluster_i
        color = cm.nipy_spectral(float(i) / n_clusters)

        ax1.fill_betweenx(np.arange(y_lower, y_upper),
                           0, ith_cluster_silhouette_values,
                           facecolor=color, edgecolor=color,
                           alpha=0.7)

        # Label the silhouette plots with their cluster numbers at the
        # middle

        ax1.text(-0.05, y_lower + 0.5 * size_cluster_i, str(i))

```

```

# Compute the new y_lower for next plot

y_lower = y_upper + 10 # 10 for the 0 samples
ax1.set_title("The silhouette plot for the various clusters.")
ax1.set_xlabel("The silhouette coefficient values")
ax1.set_ylabel("Cluster label")

# The vertical line for average silhouette score of all the
# values

ax1.axvline(x=silhouette_avg, color="red", linestyle="--")
ax1.set_yticks([]) # Clear the yaxis labels / ticks
ax1.set_xticks([-0.1, 0, 0.2, 0.4, 0.6, 0.8, 1])

# 2nd Plot showing the actual clusters formed
colors = cm.nipy_spectral(cluster_labels.astype(float) / n_clusters)
ax2.scatter(X[:, 0],
            X[:, 1],
            marker='.',
            s=30,
            lw=0,
            alpha=0.7,
            c=colors,
            edgecolor='k')

# Labeling the clusters by centers
centers = clusterer.labels_

# Time to run
print("--- %s seconds ---" % (time.time() - start_time))

#####
# Sales
#####
print('|||||')
print('|||||')
print('sales')
print('|||||')
print('|||||')
X = pca_data_sales
min_max_scaler = preprocessing.MinMaxScaler()
x_scaled = min_max_scaler.fit_transform(X)
X = x_scaled
do_cluster_analysis('sales')

#####

```

```

# Water
#####
print('|||||')
print('|||||')
print('water')
print('|||||')
print('|||||')
X = pca_data_water
min_max_scaler = preprocessing.MinMaxScaler()
x_scaled = min_max_scaler.fit_transform(X)
X = x_scaled
do_cluster_analysis('water')

#####
# Seoul
#####
print('|||||')
print('|||||')
print('seoul')
print('|||||')
print('|||||')
X = pca_data_seoul
min_max_scaler = preprocessing.MinMaxScaler()
x_scaled = min_max_scaler.fit_transform(X)
X = x_scaled
do_cluster_analysis('seoul')

```

```

|||||
||||
|||||
|||||
sales
|||||
||||
|||||
||||
=====
====
n_clusters = 2
=====
====
For n_clusters = 2 The average silhouette_score is : 0.35714256858302806
For n_clusters = 2 The average SSE is : 27.772162712889894
--- 0.346055269241333 seconds ---
=====
====
n_clusters = 3

```

```





=====
====
For n_clusters = 3 The average silhouette_score is : 0.35714256858302806
For n_clusters = 3 The average SSE is : 27.772162712889894
--- 0.23154139518737793 seconds ---
=====
====
n_clusters = 4
=====
====
For n_clusters = 4 The average silhouette_score is : 0.35714256858302806
For n_clusters = 4 The average SSE is : nan
--- 0.17214298248291016 seconds ---
=====
====
n_clusters = 5
=====
====
For n_clusters = 5 The average silhouette_score is : 0.35714256858302806
For n_clusters = 5 The average SSE is : nan
--- 0.14641571044921875 seconds ---
=====
====
n_clusters = 6
=====
====
For n_clusters = 6 The average silhouette_score is : 0.35714256858302806
For n_clusters = 6 The average SSE is : nan
--- 0.1602458953857422 seconds ---
|||||
||||
|||||
||||
water
|||||
||||
|||||
||||
=====
====
n_clusters = 2
=====
====
For n_clusters = 2 The average silhouette_score is : 0.01814249370553166
DBSCAN EXCEPTION
|||||
||||
|||||


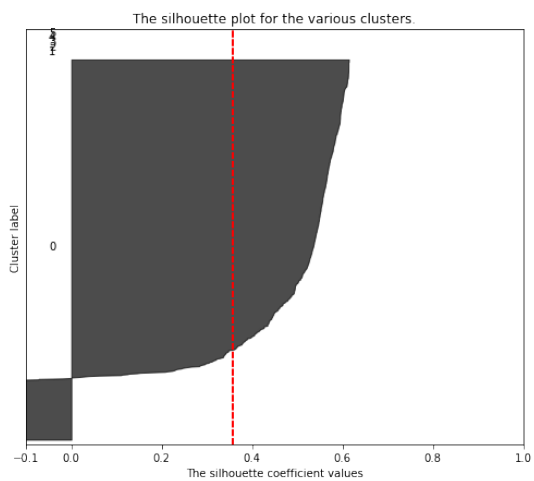



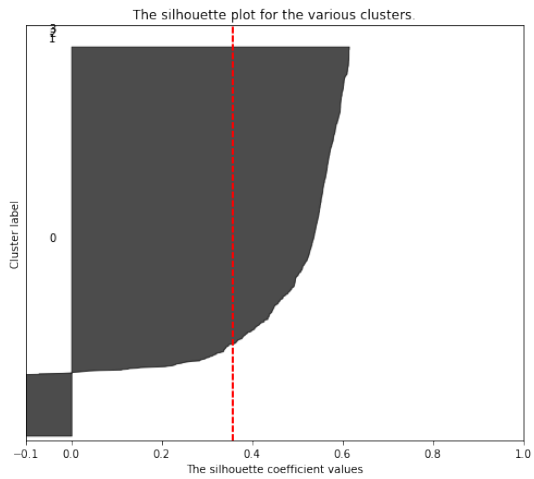
```





```

||||
seoul
|||||
|||||
|||||
=====
=====
n_clusters = 2
=====
=====
For n_clusters = 2 The average silhouette_score is : -0.48965343976481296
DBSCAN EXCEPTION

```







[]: