

19075153_O'Leary_PartB

October 24, 2021

```
[1]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import warnings
from pandas.plotting import scatter_matrix
import seaborn as sns
from sklearn.model_selection import cross_val_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report, confusion_matrix
from sklearn import tree
from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import chi2
from numpy import set_printoptions
from sklearn.decomposition import PCA
from sklearn import preprocessing
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import RANSACRegressor
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from scipy.stats import pearsonr
import sklearn.metrics as metrics

warnings.filterwarnings('ignore')
```

1 Data load and pre-processing

```
[2]: min_max_scaler = preprocessing.MinMaxScaler()

#####
# load Penrose data
#####
path = "/mnt/c/Users/T828808/Study/AUT/COMP809/Ass2/
↳Penrose_Hourly_AggregateData_Jan2016Dec2020.csv"
rawdata = pd.read_csv(path)
```

```

# Rearrange Timestamp column
# Remove incorrect values or outliers
rawdata = rawdata[rawdata['Relative Humidity (%)'] <= 100]
rawdata = rawdata[rawdata['Air Temp (C)'] <= 40]
rawdata = rawdata.dropna(how='any',axis=0)
rawdata['Timestamp (UTC+12:00)'] = pd.to_datetime(rawdata['Timestamp (UTC+12:
→00)'])
rawdata = rawdata.set_index('Timestamp (UTC+12:00)')
rawdata['Hour'] = rawdata.index.hour
rawdata['Day'] = rawdata.index.day
rawdata['Month'] = rawdata.index.month
rawdata['Year'] = rawdata.index.year

# categorise everything and create array
list_of_columns = rawdata.columns
rawdata[list_of_columns] = rawdata[list_of_columns].apply(lambda col:pd.
→Categorical(col).codes)

```

2 Inspecting outliers

```

[3]: # A function that will plot a scatterplot of the training samples and add the
→regression line
def lin_reg_plot(x_name):

    # Create the model
    slr = LinearRegression()
    slr.fit(X, y)

    # Regression metrics
    explained_variance=metrics.explained_variance_score(y, slr.predict(X))
    mean_absolute_error=metrics.mean_absolute_error(y, slr.predict(X))
    mse=metrics.mean_squared_error(y, slr.predict(X))
    median_absolute_error=metrics.median_absolute_error(y, slr.predict(X))
    r2=metrics.r2_score(y, slr.predict(X))
    corr, _ = pearsonr(rawdata[x_name], y)

    # Print stats
    print('===== ' + x_name + '
→=====')
    print('Slope: %.3f' % slr.coef_[0])
    print('Intercept: %.3f' % slr.intercept_)
    print('explained_variance: ', round(explained_variance,4))
    print('r2: ', round(r2,4))
    print('MAE: ', round(mean_absolute_error,4))
    print('MSE: ', round(mse,4))
    print('RMSE: ', round(np.sqrt(mse),4))

```

```
print('Pearsons correlation: ', round(corr,4)) # Over 0.5 or less than -0.5
↳signifies strongest correlation
```

```
# Chart
```

```
plt.scatter(X, y, c='steelblue', edgecolor='white', s=70)
plt.plot(X, slr.predict(X), color='black', lw=2)
plt.xlabel(x_name)
plt.ylabel('PM2.5 ( g/m )')
plt.show()
```

```
# Starting point X and y
```


```
X = rawdata[['Hour','Air Temp ( C)','Relative Humidity (%)','Solar Radiation (W/
↳m)','Wind Direction ( )','Wind Speed (m/s)','NO ( g/m )','NO2 ( g/m )','SO2 ( g/
↳m )']].values
y = rawdata['PM2.5 ( g/m )'].values
```

```
# Display data in 2D
```

```
x_name = 'SO2 ( g/m )'
X = rawdata[[x_name]].values
lin_reg_plot(x_name)
x_name = 'NO2 ( g/m )'
X = rawdata[[x_name]].values
lin_reg_plot(x_name)
x_name = 'NO ( g/m )'
X = rawdata[[x_name]].values
lin_reg_plot(x_name)
x_name = 'Wind Speed (m/s)'
X = rawdata[[x_name]].values
lin_reg_plot(x_name)
x_name = 'Wind Direction ( )'
X = rawdata[[x_name]].values
lin_reg_plot(x_name)
x_name = 'Solar Radiation (W/m )'
X = rawdata[[x_name]].values
lin_reg_plot(x_name)
x_name = 'Relative Humidity (%)'
X = rawdata[[x_name]].values
lin_reg_plot(x_name)
x_name = 'Air Temp ( C )'
X = rawdata[[x_name]].values
lin_reg_plot(x_name)
x_name = 'Hour'
X = rawdata[[x_name]].values
lin_reg_plot(x_name)
```


```
===== SO2 ( g/m )
=====
```

Slope: 1.011
 Intercept: 106.215
 explained_variance: 0.1236
 r2: 0.1236
 MAE: 36.4505
 MSE: 2458.9511
 RMSE: 49.5878
 Pearsons correlation: 0.3515




===== NO2 (g/m)
 =====

Slope: 0.128
 Intercept: 98.190
 explained_variance: 0.1137
 r2: 0.1137
 MAE: 36.438
 MSE: 2486.5994
 RMSE: 49.8658
 Pearsons correlation: 0.3372




```
===== NO ( g/m )
=====
```

```
Slope: 0.113
Intercept: 103.757
explained_variance: 0.1904
r2: 0.1904
MAE: 35.2568
MSE: 2271.5162
RMSE: 47.6604
Pearsons correlation: 0.4363
```




```
===== Wind Speed (m/s)
=====
Slope: -0.233
Intercept: 127.722
explained_variance: 0.0045
r2: 0.0045
MAE: 38.6166
MSE: 2792.8766
RMSE: 52.8477
Pearsons correlation: -0.0673
```



===== Wind Direction ()

=====


Slope: 0.085
 Intercept: 106.028
 explained_variance: 0.0211
 r2: 0.0211
 MAE: 38.016
 MSE: 2746.4533
 RMSE: 52.4066
 Pearsons correlation: 0.1452



```

===== Solar Radiation (W/m )
=====
Slope: 0.001
Intercept: 121.037
explained_variance: 0.0004
r2: 0.0004
MAE: 38.47
MSE: 2804.466
RMSE: 52.9572
Pearsons correlation: 0.02


```

```

===== Relative Humidity (%)
=====
Slope: -0.003
Intercept: 123.050
explained_variance: 0.0001
r2: 0.0001
MAE: 38.4961
MSE: 2805.3804
RMSE: 52.9658
Pearsons correlation: -0.0086

```



===== Air Temp (C)

=====

Slope: -3.581

Intercept: 172.655

explained_variance: 0.0722


r2: 0.0722

MAE: 37.9256

MSE: 2603.0704

RMSE: 51.0203


Pearsons correlation: -0.2687



```


===== Hour =====
Slope: -0.318
Intercept: 125.667
explained_variance:  0.0017
r2:  0.0017
MAE:  38.4984
MSE:  2800.7835
RMSE:  52.9224
Pearsons correlation:  -0.0414

```



3 Feature importance

```
[4]: # Run PCA
pca = PCA()
pca_fit = pca.fit(rawdata)
# summarize components
feat_importances = pd.Series(pca_fit.explained_variance_ratio_, index=rawdata.
    ↪columns)
feat_importances.nlargest(20).plot(kind='barh')
# Get only the first two components as they explain almost all of the variance
pca = PCA(n_components=2)
pca_fit = pca.fit(rawdata)
```



4 Summary stats for PM2.5 (g/m)

```
[5]: rawdata['PM2.5 ( g/m )'].describe()
```

```
[5]: count      10313.000000
     mean        121.846310
     std         52.970365
     min          0.000000
     25%          88.000000
     50%         113.000000
     75%         145.000000
     max         397.000000
     Name: PM2.5 ( g/m ), dtype: float64
```

5 Do the train/test split

```
[9]: # Get final X and y - we want to keep the top five attributes based on Pearsons
     ↪ correlation
     # Five topmost correlated varaiaables are 'Air Temp (C)', 'Wind Direction
     ↪ ( )', 'NO ( g/m )', 'NO2 ( g/m )', 'SO2 ( g/m )'
     X = rawdata[['Air Temp (C)', 'Wind Direction ( )', 'NO ( g/m )', 'NO2 ( g/m )', 'SO2
     ↪ ( g/m )']].values
     y = rawdata['PM2.5 ( g/m )'].values
```

```

# We need to reduce the number of categories for y - max is 397, we can
↳ probably start with 7 categories
# In this case we get about 0.45 accuracy score
# If we reduce the categories to 4, we get about 0.6 accuracy score
y = np.digitize(y,bins=[50,100,150,200,250,300,350])

# Do the train/test split
pred_train, pred_test, tar_train, tar_test = train_test_split(X, y, test_size=.
↳ 3, random_state=4)

```

6 MLP classifier

```

[10]: import pandas
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.neural_network import MLPClassifier
import matplotlib.pyplot as plt

#A function to see some of the attributes of NN
def NN_properties(model):
    loss_values = model.loss_
    print("Loss", loss_values)
    iterations = model.n_iter_
    print("iterations",iterations)
    classes_assigned= model.classes_
    print("Assigned classes", classes_assigned)

#Displaying loss curve using loss_curve method.Note that this only works with
↳ the MLP default solver "adam"
def make_plots_default(model):
    plt.plot(model.loss_curve_)
    plt.title('Loss Curve')
    plt.xlabel('Epochs')
    plt.ylabel('Loss')
    plt.show()

#A generic function to display training loss and testing accuracy of
↳ MLPClassifier
def make_plots_all(mlp, target_train, target_test,
↳ predictors_test,predictors_train):
    max_iter = 100
    accuracy = []
    losses = []
    for i in range(max_iter):
        mlp.fit(predictors_train, target_train)

```

```


        iter_acc = mlp.score(predictors_test, target_test)
        accuracy.append(iter_acc)
        losses.append(mlp.loss_)
plt.plot(accuracy, label='Test accuracy')
plt.plot(losses, label='Loss')
plt.title("Accuracy and Loss over Iterations", fontsize=14)
plt.xlabel('Iterations')
plt.legend(loc='upper right')
plt.show()

#A function for model building and calculating accuracy
def get_accuracy(target_train, target_test, predictors_test, predictors_train):
    # Two hidden layers with 10 and 5 neurons - NN
    clf = MLPClassifier(hidden_layer_sizes=(25), max_iter=100)
    #Calling the make_plots_allfunction with unfitted model
    make_plots_all(clf, target_train, target_test, predictors_test,
    ↪predictors_train)
    clf.fit(predictors_train, np.ravel(target_train, order='C'))
    predictions = clf.predict(predictors_test)
    NN_properties(clf) ##Calling NN_properties to see the model attributes
    make_plots_default(clf) ##Calling make_plots function to see the error plots
    return accuracy_score(target_test, predictions)


#train-test split
pred_train, pred_test, tar_train, tar_test = train_test_split(X, y, test_size=.
    ↪3, random_state=4)

#Calling get_accuracy function which also invoke other functions NN_properties,
    ↪make_plots, make_plots_all
print("Accuracy score: %.2f" % get_accuracy(tar_train, tar_test, pred_test,
    ↪pred_train))

```



Loss 1.4461270468227754
iterations 78
Assigned classes [0 1 2 3 4 5 6 7]



Accuracy score: 0.44

7 2-Layer MLP

```
[11]: #train-test split
pred_train, pred_test, tar_train, tar_test = train_test_split(X, y, test_size=.
    ↪3, random_state=4)

def two_layer_mlp(num_l1):
    tlmlp = MLPClassifier(hidden_layer_sizes=(num_l1, 25-num_l1,), max_iter=84)
    fit_tlmlp = tlmlp.fit(pred_train, np.ravel(tar_train, order='C'))
    pred_tlmlp = tlmlp.predict(pred_test)
    prob_tlmlp = tlmlp.predict_proba(pred_test)
    accuracy_tlmlp = accuracy_score(tar_test, pred_tlmlp)
    return accuracy_tlmlp

results = pd.DataFrame(columns=["Combination of neurons", "Accuracy"])
for i in range(1,25):
    new_row = {"Combination of neurons": str(i)+"-"+str(25-i), "Accuracy":
    ↪str(two_layer_mlp(i))}
    results = results.append(new_row, ignore_index=True)

# Drop results out to a table
```

```
results
```

```
[11]: Combination of neurons      Accuracy
0      1,24 0.43956043956043955
1      2,23 0.38978668390433097
2      3,22 0.4188752424046542
3      4,21 0.41047188106011634
4      5,20 0.3784744667097608
5      6,19 0.4460245636716225
6      7,18 0.43956043956043955
7      8,17 0.4253393665158371
8      9,16 0.4414996767937944
9     10,15 0.4411764705882353
10    11,14 0.4434389140271493
11    12,13 0.4137039431157078
12    13,12 0.4308338720103426
13    14,11 0.43244990303813835
14    15,10 0.4114414996767938
15    16,9  0.44505494505494503
16    17,8  0.4502262443438914
17    18,7  0.4204912734324499
18    19,6  0.44182288299935357
19    20,5  0.4175824175824176
20    21,4  0.41952165481577247
21    22,3  0.4230769230769231
22    23,2  0.4191984486102133
23    24,1  0.4191984486102133
```

```
[ ]:
```