

COMP809 – Data Mining and Machine Learning

Lab 7 – Clustering with Kmeans

This lab will cover Kmeans clustering.

Study the code provided below and run it in Python. The code uses SSE (it is referred to as distortion in the code) and the cluster silhouette measure for evaluation.

The code also compares two different cluster configurations, one with $K=3$ and $K=2$. You will observe that $K=3$ produces a better cluster configuration.

Extend the code to use the Iris dataset. Once it is working use PCA as a pre processing step prior to applying Kmeans.

KMeans module from sklearn is used to perform the clustering
#Go to <https://scikit-learn.org/stable/modules/generated/sklearn.cluster.KMeans.html>
for more details.

```
from IPython.display import Image
from sklearn.datasets import make_blobs
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
import numpy as np
from matplotlib import cm
from sklearn.metrics import silhouette_samples
```

#Create a sample data set using make_blobs. This particular dataset has 2 features and 3 clusters.

```
X, y = make_blobs(n_samples=150,
                  n_features=2,
                  centers=3,
                  cluster_std=0.5,
                  shuffle=True,
                  random_state=0)
```

#Display the dataset

```
plt.scatter(X[:, 0], X[:, 1], c='white', marker='o',
            edgecolor='black', s=50)
plt.grid()
plt.tight_layout()
```

#Customize the path to save images

```
plt.savefig('P:/COMP809/11_01.png', dpi=300)
plt.show()
```

#Apply k-means clustering with 3 centroids

```
km = KMeans(n_clusters=3,
            init='random',
            n_init=10,
            max_iter=300,
            tol=1e-04,
            random_state=0)
```

#We set n_init=10 to run the k-means clustering algorithms 10 times independently with different random centroids to choose the final model as the one with the lowest SSE.

#Predicting cluster labels

```
y_km = km.fit_predict(X)
```

#Visualize the clusters identified(using y_km) together with cluster labels.

```
plt.scatter(X[y_km == 0, 0],
            X[y_km == 0, 1],
            s=50, c='lightgreen',
            marker='s', edgecolor='black',
            label='cluster 1')

plt.scatter(X[y_km == 1, 0],
            X[y_km == 1, 1],
            s=50, c='orange',
            marker='o', edgecolor='black',
            label='cluster 2')

plt.scatter(X[y_km == 2, 0],
            X[y_km == 2, 1],
            s=50, c='lightblue',
            marker='v', edgecolor='black',
            label='cluster 3')

plt.scatter(km.cluster_centers_[0, 0],
            km.cluster_centers_[0, 1],
            s=250, marker='*',
            c='red', edgecolor='black',
            label='centroids')

plt.legend(scatterpoints=1)
plt.grid()
plt.tight_layout()
plt.savefig('P:/COMP809/11_02.png', dpi=300)
plt.show()

#Calculating Distortion
print('Distortion: %.2f' % km.inertia_)
distortions = []
```

#Observing the behaviour of the distortion with the number of clusters.

#Using elbow method to find optimum number of clusters

```
for i in range(1, 11):
    km = KMeans(n_clusters=i,
                init='k-means++',
                n_init=10,
                max_iter=300,
                random_state=0)

    km.fit(X)
    distortions.append(km.inertia_)

plt.plot(range(1, 11), distortions, marker='o')
plt.xlabel('Number of clusters')
plt.ylabel('Distortion')
plt.tight_layout()
plt.savefig('P:/COMP809/11_03.png', dpi=300)
plt.show()
```

#Quantifying the quality of clustering via silhouette plots.

#k-means++ gives better clustering/performance than classic approach(init='random'). Always recommended to use k-means++

```
km = KMeans(n_clusters=3,
            init='k-means++',
            n_init=10,
            max_iter=300,
            tol=1e-04,
            random_state=0)
y_km = km.fit_predict(X)
```

#Silhouette Measure/plots for cluster evaluation. Higher the score better the clustering.

```
cluster_labels = np.unique(y_km)
n_clusters = cluster_labels.shape[0]
silhouette_vals = silhouette_samples(X, y_km,
metric='euclidean')
y_ax_lower, y_ax_upper = 0, 0
yticks = []
for i, c in enumerate(cluster_labels):
```

```

    c_silhouette_vals = silhouette_vals[y_km == c]
    c_silhouette_vals.sort()
    y_ax_upper += len(c_silhouette_vals)
    color = cm.jet(float(i) / n_clusters)
    plt.barh(range(y_ax_lower, y_ax_upper), c_silhouette_vals,
height=1.0,edgecolor='none', color=color)

    yticks.append((y_ax_lower + y_ax_upper) / 2.)
    y_ax_lower += len(c_silhouette_vals)

silhouette_avg = np.mean(silhouette_vals)
plt.axvline(silhouette_avg, color="red", linestyle="--")

plt.yticks(yticks, cluster_labels + 1)
plt.ylabel('Cluster')
plt.xlabel('Silhouette coefficient')

plt.tight_layout()
plt.savefig('P:/COMP809/11_04.png', dpi=300)
plt.show()

```

Comparison to "bad" clustering: Same data samples are clustered using 2 cluters.

```

km = KMeans(n_clusters=2,
            init='k-means++',
            n_init=10,
            max_iter=300,
            tol=1e-04,
            random_state=0)
y_km = km.fit_predict(X)

plt.scatter(X[y_km == 0, 0],
            X[y_km == 0, 1],
            s=50,
            c='lightgreen',
            edgecolor='black',
            marker='s',
            label='cluster 1')

plt.scatter(X[y_km == 1, 0],
            X[y_km == 1, 1],
            s=50,
            c='orange',

```

```

        edgecolor='black',
        marker='o',
        label='cluster 2')

plt.scatter(km.cluster_centers_[ :, 0], km.cluster_centers_[ :,
1], s=250, marker='*', c='red', label='centroids')
plt.legend()
plt.grid()
plt.tight_layout()
plt.savefig('P:/COMP809/11_05.png', dpi=300)
plt.show()

cluster_labels = np.unique(y_km)
n_clusters = cluster_labels.shape[0]
silhouette_vals = silhouette_samples(X, y_km,
metric='euclidean')
y_ax_lower, y_ax_upper = 0, 0
yticks = []
for i, c in enumerate(cluster_labels):
    c_silhouette_vals = silhouette_vals[y_km == c]
    c_silhouette_vals.sort()
    y_ax_upper += len(c_silhouette_vals)
    color = cm.jet(float(i) / n_clusters)
    plt.barh(range(y_ax_lower, y_ax_upper), c_silhouette_vals,
height=1.0,
            edgecolor='none', color=color)

    yticks.append((y_ax_lower + y_ax_upper) / 2.)
    y_ax_lower += len(c_silhouette_vals)

silhouette_avg = np.mean(silhouette_vals)
plt.axvline(silhouette_avg, color="red", linestyle="--")

plt.yticks(yticks, cluster_labels + 1)
plt.ylabel('Cluster')
plt.xlabel('Silhouette coefficient')

plt.tight_layout()
plt.savefig('P:/COMP809/11_06.png', dpi=300)
plt.show()

```