





DATA MINING AND MACHINE LEARNING

Clustering Algorithms 3

Hierarchical Clustering

- Produces a set of nested clusters organized as a hierarchical tree
- Can be visualized as a dendrogram
 - *A tree like diagram that records the sequences of merges or splits*



Hierarchical Clustering


- ▶ Two main types of hierarchical clustering
 - *Agglomerative:*
 - Start with the **points as individual clusters**
 - At each step, merge the closest pair of clusters until only one cluster (or k clusters) left
 - *Divisive:*
 - Start with **one, all-inclusive cluster**
 - At each step, split a cluster until each cluster contains a point (or there are k clusters)
- ▶ Traditional hierarchical algorithms use a similarity or distance matrix
 - *Merge or split one cluster at a time*

Agglomerative Clustering Algorithm

- More popular hierarchical clustering technique
- Basic algorithm is straightforward
 - Compute the proximity matrix*
 - Let each data point be a cluster*
 - Repeat**
 - Merge the two closest clusters
 - Update the proximity matrix
 - Until** only a single cluster remains
- Key operation is the computation of the proximity of two clusters
 - *Different approaches to defining the distance between clusters distinguish the different algorithms*

Starting Situation

- Start with clusters of individual points and a proximity matrix




	p1	p2	p3	p4	p5	...
p1						
p2						
p3						
p4						
p5						
.						
.						
.						

Proximity Matrix

p1 p2 p3 p4 ... p9 p10 p11 p12


Intermediate Situation

- After some merging steps, we have some clusters




	C1	C2	C3	C4	C5
C1					
C2					
C3					
C4					
C5					

Proximity Matrix




Intermediate Situation

- We want to merge the two closest clusters (C2 and C5) and update the proximity matrix.




	C1	C2	C3	C4	C5
C1					
C2					
C3					
C4					
C5					

Proximity Matrix



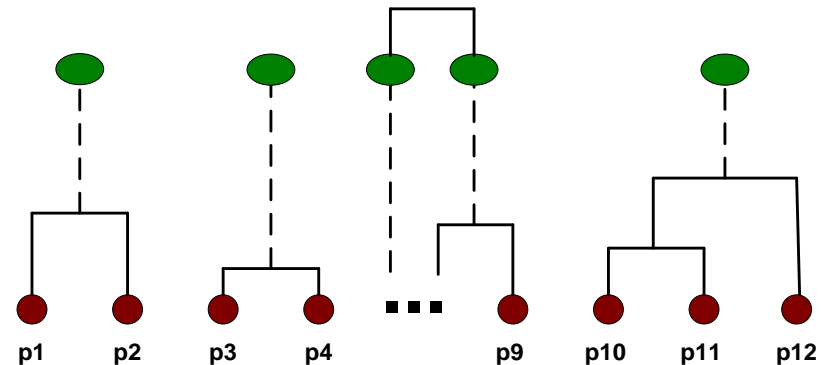
After Merging

- The question is “How do we update the proximity matrix?”




		$C2 \cup C5$			
		C1	C5	C3	C4
C1			?		
$C2 \cup C5$?		?	?
C3			?		
C4			?		

Proximity Matrix



How to Define Inter-Cluster Similarity




- MIN
- MAX
- Group Average
- Distance Between Centroids

	p1	p2	p3	p4	p5	...
p1						
p2						
p3						
p4						
p5						
.						
.						
.						

Proximity Matrix

How to Define Inter-Cluster Similarity




- MIN
- MAX
- Group Average
- Distance Between Centroids

	p1	p2	p3	p4	p5	...
p1						
p2						
p3						
p4						
p5						
.						
.						
.						

Proximity Matrix

How to Define Inter-Cluster Similarity




- MIN
- MAX
- Group Average
- Distance Between Centroids

	p1	p2	p3	p4	p5	...
p1						
p2						
p3						
p4						
p5						
.						
.						
.						

· **Proximity Matrix**

How to Define Inter-Cluster Similarity




- MIN
- MAX
- Group Average
- Distance Between Centroids

	p1	p2	p3	p4	p5	...
p1						
p2						
p3						
p4						
p5						
.						
.						
.						

Proximity Matrix

How to Define Inter-Cluster Similarity



- MIN
- MAX
- Group Average
- Distance Between Centroids


	p1	p2	p3	p4	p5	...
p1						
p2						
p3						
p4						
p5						
.						
.						

· **Proximity Matrix**


Cluster Similarity: MIN or Single Link

- Similarity of two clusters is based on the two most similar (closest) points in the different clusters
 - *Determined by one pair of points, i.e., by one link in the proximity graph.*


	I1	I2	I3	I4	I5
I1	1.00	0.90	0.10	0.65	0.20
I2	0.90	1.00	0.70	0.60	0.50
I3	0.10	0.70	1.00	0.40	0.30
I4	0.65	0.60	0.40	1.00	0.80
I5	0.20	0.50	0.30	0.80	1.00



Hierarchical Clustering: MIN




Nested Clusters



Dendrogram

Dendrogram: Hierarchical Clustering

- Clustering obtained by cutting the dendrogram at a desired level: each **connected** component forms a cluster.



Python support for Hierarchical Clustering

```
class sklearn.cluster.AgglomerativeClustering(n_clusters=2, *,  
affinity='euclidean',memory=None, connectivity=None,  
compute_full_tree='auto', linkage='ward', distance_threshold=None)
```

Tuneable parameters are

- **n_clusters** – in hierarchical clustering need to specify a cut-of point
- **linkage** – how the distance between points is defined
- **distance_threshold** – neighboring clusters with greater value than this will not be merged; by controlling this you effectively set the height of the tree

Strengths of Hierarchical Clustering



- Do not have to assume any particular number of clusters
 - *Any desired number of clusters can be obtained by ‘cutting’ the dendrogram at the proper level*
- They may correspond to meaningful taxonomies
 - *Example in biological sciences (e.g., animal kingdom, phylogeny reconstruction, ...)*

Density-Based Clustering Methods

- Clustering based on density (local cluster criterion), such as density-connected points or based on an explicitly constructed density function
- Major features:
 - *Discover clusters of arbitrary shape*
 - *Handle noise*
 - *One scan*
 - *Need density parameters*
- Several interesting studies:
 - DBSCAN: Ester, et al. (KDD '96)
 - DENCLUE: Hinneburg & D. Keim (KDD '98/2006)
 - OPTICS: Ankerst, et al (SIGMOD '99).
 - CLIQUE: Agrawal, et al. (SIGMOD '98)

Density Based Clustering

- DBSCAN (Density Based Spatial Clustering of Applications with Noise) is a density based scheme that works well with non spherical clusters.




DBSCAN Algorithm

(<http://www2.cs.uh.edu/~ceick/7363/Papers/dbscan.pdf>)

- Density = number of points within a specified radius r (Eps)
- A point is a **core point** if it has more than a specified number of points (MinPts) within Eps
 - *These are points that are at the interior of a cluster*
- A **border point** has fewer than MinPts within Eps, but is in the neighborhood of a core point
 - *Greater **MinPts** and smaller r => Higher Density*
- A **noise point** is any point that is not a core point or a border point.

DBSCAN: Core, Border, and Noise Points




DBSCAN

DBSCAN uses these two parameters (**Eps** and **MinPts**) to divide the sample space into 3 different categories:

- **Core points**: a sample **s** is a core point IF at least MinPts points are within distance **Eps** of it (including **s**).
- **Reachable (border) points**: a sample **t** is reachable from **s** IF point **t** is within distance **Eps** from core point **s**.
- **Noise points**: All samples not reachable from any other sample is said to be noise points.

DBSCAN



This image extracted from:

Schubert, Erich; Sander, Jörg; Ester, Martin; [Kriegel, Hans Peter](#); Xu, Xiaowei (July 2017). ["DBSCAN Revisited, Revisited: Why and How You Should \(Still\) Use DBSCAN"](#). *ACM Trans. Database Syst.* **42** (3): 19:1–19:21.

DBSCAN Demo

<https://www.youtube.com/watch?v=h53WMllmUuc>

Outline Algorithm

1. Create a cluster (C) for each set of core points (S)
2. Include in C any points reachable from any core point in set S.

DBSCAN: Detailed Algorithm

ALGORITHM 1: Pseudocode of Original Sequential DBSCAN Algorithm

```
Input: DB: Database
Input:  $\varepsilon$ : Radius
Input: minPts: Density threshold
Input: dist: Distance function
Data: label: Point labels, initially undefined

1 foreach point p in database DB do                                // Iterate over every point
2   if label(p)  $\neq$  undefined then continue                        // Skip processed points
3   Neighbors N  $\leftarrow$  RANGEQUERY(DB, dist, p,  $\varepsilon$ )           // Find initial neighbors
4   if  $|N| < \text{minPts}$  then                                           // Non-core points are noise
5     label(p)  $\leftarrow$  Noise
6     continue
7   c  $\leftarrow$  next cluster label                                     // Start a new cluster
8   label(p)  $\leftarrow$  c
9   Seed set S  $\leftarrow N \setminus \{p\}$                              // Expand neighborhood
10  foreach q in S do
11    if label(q) = Noise then label(q)  $\leftarrow$  c
12    if label(q)  $\neq$  undefined then continue
13    Neighbors N  $\leftarrow$  RANGEQUERY(DB, dist, q,  $\varepsilon$ )
14    label(q)  $\leftarrow$  c
15    if  $|N| < \text{minPts}$  then continue                                // Core-point check
16    S  $\leftarrow S \cup N$ 
```

- Above code taken from Schubert et al.

DBSCAN Algorithm (simplified)

1. Create a graph whose nodes are the points to be clustered
2. For each **core-point** c create an edge from c to every point p in the ε -neighborhood of c
3. Set N to the nodes of the graph;
4. If N does not contain any core points terminate
5. Pick a core point c in N
6. Let X be the set of nodes that can be reached from c by going forward;
 1. *create a cluster containing $X \cup \{c\}$*
 2. $N = N / (X \cup \{c\})$
7. Continue with step 4

Note: points that are not assigned to any cluster are outliers;


Parameter estimation

- We first try to establish a value for *min_pts* and then set a value for ϵ
- As a rule of thumb we set $\text{min_pts} \geq 2 \times D$ where D is the number of dimensions
- For datasets which are large or are noisy values of $\text{min_pts} > 2 \times D$ are appropriate
- Now use the *min_pts* parameter and vary the values of ϵ until the silhouette measure is maximized.

DBSCAN: Core, Border and Noise Points




Original Points




Point types: **core**,
border and **noise**

Eps = 10, MinPts = 4

When DBSCAN Works Well




Original Points



Clusters


- Resistant to Noise
- Can handle clusters of different shapes and sizes

When DBSCAN Does NOT Work Well




Original Points

- Varying densities
- High-dimensional data



(MinPts=4, Eps=9.75).



(MinPts=4, Eps=9.92)

Advantages

- Automatically finds the number of clusters once the ϵ and min_pts parameters are specified
- Can capture non spherical clusters
- Can capture noise points in dataset

Disadvantages

- Can be slow for large datasets as complexity is $O(n^2)$
- Can produce poor quality solutions when data contains clusters of widely different density
- High dimensional datasets can also result in poor clustering as distances of points to clusters get distorted in space (affects other clustering algorithms as well)

Python support for DBSCAN

- Supported by `sklearn.dbSCAN`

```
class sklearn.cluster.DBSCAN(eps=0.5, *, min_samples=5,  
metric='euclidean', metric_params=None,  
algorithm='auto', leaf_size=30, p=None, n_jobs=None)
```

- Tuneable parameters are `eps` and `min_samples` (`min_pts`)
- Returns *labels*, a one dimensional array of cluster indexes that samples belongs to; noisy samples are flagged by -1

Other Clustering Algorithms

Many other widely used algorithms exist such as:

- SOM (Self Organizing Map)
- EM (Expectation Maximization)
- BIRCH(Balanced Iterative Reducing and Clustering)
- Cascade K means
- etc.