




DATA MINING & MACHINE LEARNING

Convolutional Neural Networks




Fully Connected Neural Network

- We know it is good to learn a small model.
- From this fully connected model, do we really need all the edges?
- Can some of these be shared?




Problems using FC Layers on Images

- How to process a tiny image with FC layers




Problems using FC Layers on Images

- How to process a tiny image with FC layers




Problems using FC Layers on Images

- How to process a tiny image with FC layers




Problems using FC Layers on Images

- How to process a tiny image with FC layers




Problems using FC Layers on Images

- How to process a normal image with FC layers



Problems using FC Layers on Images

- How to process a normal image with FC layers



Why not simply more FC Layers?

We cannot make networks arbitrarily complex

- Why not just go deeper and get better?
 - No structure!!
 - It is just brute force!
 - Optimization becomes hard
 - Performance plateaus / drops!

Better Way than FC ?

- We want to restrict the degrees of freedom
 - We want a layer with structure
 - Weight sharing → using the same weights for different parts of the image

Using CNNs in Computer Vision

Classification




CAT

**Classification
+ Localization**



CAT

Object Detection



CAT, DOG, DUCK

**Instance
Segmentation**



CAT, DOG, DUCK

Single object

Multiple objects

Convolutional NN

In 1995, **Yann LeCun** and **Yoshua Bengio** introduced the concept of convolutional neural networks.

Convolutional Neural Networks is extension of traditional Multi-layer Perceptron, based on 3 ideas:

1. Local receive fields
2. Shared weights
3. Spatial / temporal sub-sampling

See LeCun paper (1998) on text recognition:

<http://yann.lecun.com/exdb/publis/pdf/lecun-01a.pdf>



About CNN's

- CNN's Were **neurobiologically** motivated by the findings of locally sensitive and orientation-selective nerve cells in the visual cortex.
- They designed a network structure that implicitly extracts relevant features.
- Convolutional Neural Networks are a special kind of **multi-layer neural networks**.


About CNN's

- CNN is a **feed-forward** network that can extract topological properties from an image.
- Like almost every other neural networks they **are trained** with a version of the **back-propagation algorithm**.
- Convolutional Neural Networks are designed to **recognize visual patterns** directly from pixel images with minimal preprocessing.
- They can recognize patterns with extreme variability (such as handwritten characters).

Convolutions

What are Convolutions?


Discrete case: box filter



‘Slide’ filter kernel from left to right; at each position, compute a single value in the output data

What are Convolutions?


Discrete case: box filter



$$4 * (1/3) + 3 * (1/3) + 2 * (1/3) = 3$$

What are Convolutions?


Discrete case: box filter



$$3 * (1/3) + 2 * (1/3) + (-5) * (1/3) = 0$$

What are Convolutions?

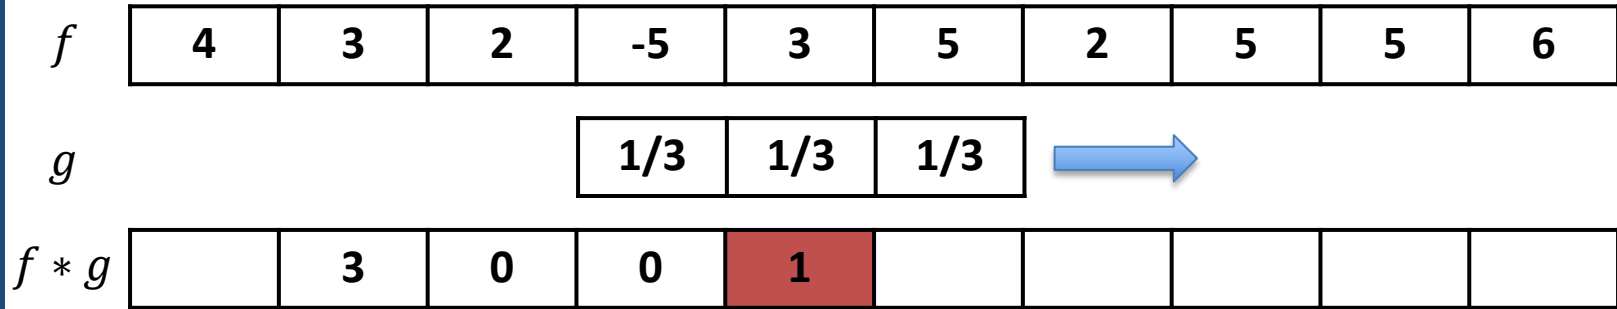
Discrete case: box filter



$$2 * (1/3) + (-5) * (1/3) + 3 * (1/3) = 0$$

What are Convolutions?

Discrete case: box filter



$$-5 * (\frac{1}{3}) + 3 * (\frac{1}{3}) + 5 * (\frac{1}{3}) = 1$$

What are Convolutions?

Discrete case: box filter

f	4	3	2	-5	3	5	2	5	5	6
g					1/3	1/3	1/3			
$f * g$		3	0	0	1	10/3				

$$3 * (1/3) + 5 * (1/3) + 2 * (1/3) = 10/3$$

What are Convolutions?


Discrete case: box filter

f	4	3	2	-5	3	5	2	5	5	6
g						1/3	1/3	1/3		
$f * g$		3	0	0	1	10/3	4			

$$5 * (1/3) + 2 * (1/3) + 5 * (1/3) = 4$$

What are Convolutions?

Discrete case: box filter


f	4	3	2	-5	3	5	2	5	5	6
g							1/3	1/3	1/3	
$f * g$		3	0	0	1	10/3	4	4		

$$2 * (1/3) + 5 * (1/3) + 5 * (1/3) = 4$$

What are Convolutions?

Discrete case: box filter

f	4	3	2	-5	3	5	2	5	5	6
g								1/3	1/3	1/3
$f * g$		3	0	0	1	10/3	4	4	16/3	



$$5 * (1/3) + 5 * (1/3) + 6 * (1/3) = 16/3$$

What are Convolutions?

Discrete case: box filter

4	3	2	-5	3	5	2	5	5	6
---	---	---	----	---	---	---	---	---	---

$\frac{1}{3}$	$\frac{1}{3}$	$\frac{1}{3}$
---------------	---------------	---------------

??	3	0	0	1	$\frac{10}{3}$	4	4	$\frac{16}{3}$??
----	---	---	---	---	----------------	---	---	----------------	----

What to do at boundaries?

What are Convolutions?

Discrete case: box filter

4	3	2	-5	3	5	2	5	5	6
---	---	---	----	---	---	---	---	---	---

$\frac{1}{3}$	$\frac{1}{3}$	$\frac{1}{3}$
---------------	---------------	---------------

??	3	0	0	1	$\frac{10}{3}$	4	4	$\frac{16}{3}$??
----	---	---	---	---	----------------	---	---	----------------	----

What to do at boundaries?

Option 1: Shrink

3	0	0	1	$\frac{10}{3}$	4	4	$\frac{16}{3}$
---	---	---	---	----------------	---	---	----------------

What are Convolutions?

Discrete case: box filter

0	4	3	2	-5	3	5	2	5	5	6	0
---	---	---	---	----	---	---	---	---	---	---	---

1/3	1/3	1/3
-----	-----	-----

??	3	0	0	1	10/3	4	4	16/3	??
----	---	---	---	---	------	---	---	------	----

$0 \cdot \frac{1}{3} + 4 \cdot \frac{1}{3} + 3 \cdot \frac{1}{3} = \frac{7}{3}$

What to do at boundaries?

Option 2: Pad (often 0's)

7/3	3	0	0	1	10/3	4	4	16/3	11/3
-----	---	---	---	---	------	---	---	------	------

Convolutions on Images

Image 5x5

-5	3	2	-5	3
4	3	2	1	-3
1	0	3	3	5
-2	0	1	4	4
5	6	7	9	-1

Kernel 3x3

0	-1	0
-1	5	-1
0	-1	0



Output 3x3

6		

$$3 * (-1) + 4 * (-1) + 3 * (5) + 2 * (-1) + 0 * (-1) \\ = 15 - 9 = 6$$

Convolutions on Images

Image 5x5

-5	3	2	-5	3
4	3	2	1	-3
1	0	3	3	5
-2	0	1	4	4
5	6	7	9	-1

Kernel 3x3

0	-1	0
-1	5	-1
0	-1	0



Output 3x3

6	1	

$$2 * (-1) + 3 * (-1) + 2 * (5) + 1 * (-1) + 3 * (-1) \\ = 10 - 9 = 1$$

Convolutions on Images

Image 5x5

-5	3	2	-5	3
4	3	2	1	-3
1	0	3	3	5
-2	0	1	4	4
5	3	7	9	1

Kernel 3x3

0	-1	0
-1	5	-1
0	-1	0



Output 3x3

6	1	8

$$\begin{aligned} & -5 * (-1) + 2 * (-1) + 1 * (5) + (-3) * (-1) + 3 * (-1) \\ & = 5 + 3 = 8 \end{aligned}$$

Convolutions on Images

Image 5x5

-5	3	2	-5	3
4	3	2	1	-3
1	0	3	3	5
-2	0	1	4	4
5	6	7	9	-1

Kernel 3x3

0	-1	0
-1	5	-1
0	-1	0



Output 3x3

6	1	8
-7		

$$3 * (-1) + 1 * (-1) + 0 * (5) + 3 * (-1) + 0 * (-1) \\ = 0 - 7 = -7$$

Convolutions on Images

Image 5x5

-5	3	2	-5	3
4	3	2	1	-3
1	0	3	3	5
-2	0	1	4	4
5	6	7	9	-1

Kernel 3x3

0	-1	0
-1	5	-1
0	-1	0



Output 3x3

6	1	8
-7	9	

$$2 * (-1) + 0 * (-1) + 3 * (5) + 3 * (-1) + 1 * (-1) \\ = 15 - 6 = 9$$

Convolutions on Images

Image 5x5

-5	3	2	-5	3
4	3	2	1	-3
1	0	3	3	5
-2	0	1	4	4
5	6	7	9	-1

Kernel 3x3

0	-1	0
-1	5	-1
0	-1	0



Output 3x3

6	1	8
-7	9	2

$$\begin{aligned} &= 1 * (-1) + 3 * (-1) + 3 * (5) + 5 * (-1) + 4 * (-1) \\ &= 15 - 13 = 2 \end{aligned}$$

Convolutions on Images

Image 5x5

-5	3	2	-5	3
4	3	2	1	-3
1	0	3	3	5
-2	0	1	4	4
5	6	7	9	-1

Kernel 3x3

0	-1	0
-1	5	-1
0	-1	0



Output 3x3

6	1	8
-7	9	2
-5		

$$0 * (-1) + (-2) * (-1) + 0 * (5) + 1 * (-1) + 6 * (-1) \\ = 2 - 7 = -5$$

Convolutions on Images

Image 5x5

-5	3	2	-5	3
4	3	2	1	-3
1	0	3	3	5
-2	0	1	4	4
5	6	7	9	-1

Kernel 3x3

0	-1	0
-1	5	-1
0	-1	0



Output 3x3

6	1	8
-7	9	2
-5	-9	

$$3 * (-1) + 0 * (-1) + 1 * (5) + 4 * (-1) + 7 * (-1) \\ = 5 - 14 = -9$$

Convolutions on Images

Image 5x5

-5	3	2	-5	3
4	3	2	1	-3
1	0	3	3	5
-2	0	1	4	4
5	6	7	9	-1

Kernel 3x3

0	-1	0
-1	5	-1
0	-1	0



Output 3x3


6	1	8
-7	9	2
-5	-9	3

$$3 * (-1) + 1 * (-1) + 4 * (5) + 4 * (-1) + 9 * (-1) = 20 - 17 = 3$$

Image Filters

- Each kernel gives us a different image filter

Input



Edge detection

$$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$$




Box mean

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



Sharpen

$$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$



Gaussian blur


$$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$

LET'S LEARN THESE FILTERS!

Convolutions on RGBImages


width height depth

image $32 \times 32 \times 3$



Depth dimension **must** match;
i.e., filter extends the full depth of the
input

filter $5 \times 5 \times 3$




Convolve filter with image
i.e., 'slide' over it and:

- apply filter at each location
- dot products

Images have depth: e.g. RGB -> 3 channels

Convolutions on RGBImages

$32 \times 32 \times 3$ image (pixels X)




1 number at a time:
equal to dot product between
filter weights w and x_i - th chunk of
the image. Here: $5 \times 5 \times 3 = 75$ -dim
dot product + bias

$$z_i = w^T x_i + b$$

$(5 \times 5 \times 3) \times 1$ $(5 \times 5 \times 3) \times 1$ 1

Convolutions on RGBImages

32×32×3 image




5×5×3 filter



Slide over all spatial locations x_i
and compute all output z_i ;
w/o padding, there are
 28×28 locations


Activation map
(also feature map)



Convolution Layer


A convolutional layer

A CNN is a neural network with some convolutional layers (and some other layers). A convolutional layer has a number of filters that does convolutional operation.



Convolution Layer

32×32×3 image




5×5×3 filter



Let's apply a different filter
with different weights!


Activation maps




Convolution Layer

Convolution “Layer”

32×32×3 image



Activation maps



Let's apply ****five**** filters,
each with different weights!


Convolution Layer

- A basic layer is defined by
 - Filter width and height (depth is implicitly given)
 - Number of different filter banks (#weightsets)
- Each filter captures a different image characteristic

Different Filters




Low-Level
Feature



- Each filter captures different image characteristics:
 - Horizontal edges
 - Vertical edges
 - Circles
 - Squares
 - ...

Dimensions of a Convolution Layer

Convolution Layers: Dimensions




Input: 7×7

Filter: 3×3

Output: 5×5

Convolution Layers: Dimensions




Input: 7×7

Filter: 3×3

Output: 5×5

Convolution Layers: Dimensions




Input: 7×7

Filter: 3×3

Output: 5×5

Convolution Layers: Dimensions




Input: 7×7

Filter: 3×3

Output: 5×5

Convolution Layers: Dimensions




Input: 7×7

Filter: 3×3

Output: 5×5

Convolution Layers: Stride

With a stride of 1



Input: 7×7

Filter: 3×3


Stride: 1

Output: 5×5

Stride of S : apply filter every S -th spatial location; i.e. subsample the image

Convolution Layers: Stride

With a stride of 2



Input: 7×7

Filter: 3×3

Stride: 2

Output: 3×3

Convolution Layers: Stride

With a stride of 2




Image 7×7

Input: 7×7

Filter: 3×3

Stride: 2

Output: 3×3

Convolution Layers: Stride

With a stride of 2




Image 7×7

Input: 7×7


Filter: 3×3

Stride: 2

Output: 3×3

Convolution Layers: Stride

With a stride of 3



Input: 7×7


Filter: 3×3

Stride: 3

Output: $? \times ?$

Convolution Layers: Stride

With a stride of 3



Input: 7×7

Filter: 3×3

Stride: 3

Output: $? \times ?$

Convolution Layers: Stride

With a stride of 3




Image 7x7

Input: 7×7
Filter: 3×3
Stride: 3
Output: $? \times ?$

Does not really fit (remainder left)
→ **Illegal stride for input & filter size!**

Convolution Layers: Dimensions

Input width of N			Filter height of F			
Filter	width	of F				

Image 7×7

Input: $N \times N$

Filter: $F \times F$

Stride: S

Output: $\left(\frac{N-F}{S} + 1\right) \times \left(\frac{N-F}{S} + 1\right)$

$$N = 7, F = 3, S = 1: \frac{7-3}{1} + 1 = 5$$


$$N = 7, F = 3, S = 2: \frac{7-3}{2} + 1 = 3$$

$$N = 7, F = 3, S = 3: \frac{7-3}{3} + 1 = 2.3$$

Fractions are illegal


Convolution Layers: Dimensions

Input Image



Shrinking down so quickly ($32 \rightarrow 28 \rightarrow 24 \rightarrow 20$) is typically not a good idea...

Convolution Layers: Padding



Why padding?

- Sizes get small too quickly
- Corner pixel is only used once

Convolution Layers: Padding

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

Image 7x7 + zero padding

Why padding?

- Sizes get small too quickly
- Corner pixel is only used once

Convolution Layers: Padding

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

Image 7×7 + zero padding

Input ($N \times N$): 7×7

Filter ($F \times F$): 3×3

Padding (P): 1

Stride (S): 1

Output 7×7



Most common is 'zero' padding

Output Size:

$$\left(\left\lfloor \frac{N + 2 * P - F}{S} \right\rfloor + 1 \right) \times \left(\left\lfloor \frac{N + 2 * P - F}{S} \right\rfloor + 1 \right)$$

$\lfloor \rfloor$ denotes the floor operator (as in practice an integer division is performed)

Convolution Layers: Padding

0	0	0	0	0	0	0	0	0
0								0
0								0
0								0
0								0
0								0
0								0
0								0
0	0	0	0	0	0	0	0	0

Image 7x7 + zero padding

Types of convolutions:

- **Valid convolution:** using no padding
- **Same convolution:** output=input size

Set padding to $P = \frac{F-1}{2}$

Convolution Layers: Dimensions

Example

Input image: $32 \times 32 \times 3$

10 filters 5×5

Stride 1


Pad 2

Depth of 3 is implicitly given

Output size is:

$$\frac{32 + 2 \cdot 2 - 5}{1} + 1 = 32$$

i.e. $32 \times 32 \times 10$



Remember

$$\text{Output: } \left(\left\lfloor \frac{N + 2 \cdot P - F}{S} \right\rfloor + 1 \right) \times \left(\left\lfloor \frac{N + 2 \cdot P - F}{S} \right\rfloor + 1 \right)$$

Convolution Layers: Dimensions

Example

Input image: $32 \times 32 \times 3$

10 filters 5×5


Stride 1

Pad 2

Output size is:

$$\frac{32 + 2 \cdot 2 - 5}{1} + 1 = 32$$

i.e. $32 \times 32 \times 10$



Remember

$$\text{Output: } \left(\left\lfloor \frac{N + 2 \cdot P - F}{s} \right\rfloor + 1 \right) \times \left(\left\lfloor \frac{N + 2 \cdot P - F}{s} \right\rfloor + 1 \right)$$

Convolution Layers: Dimensions


Example

Input image: $32 \times 32 \times 3$

10 filters 5×5

Stride 1

Pad 2



Number of parameters (weights):

Each filter has $5 \times 5 \times 3 + 1 = 76$ params (+1 for bias)


-> $76 \times 10 = 760$ parameters in layer

Example

- You are given a convolutional layer with 4 filters, kernel size 5, stride 1, and no padding that operates on an RGB image.
- Q1: What are the dimensions and the shape of its weight tensor?
 - ❑ A1: (3, 4, 5, 5)
 - ❑ A2: (4, 5, 5)
 - ❑ A3: depends on the width and height of the image

Example


- You are given a convolutional layer with 4 filters, kernel size 5, stride 1, and no padding that operates on an RGB image.
- Q1: What are the dimensions and the shape of its weight tensor?



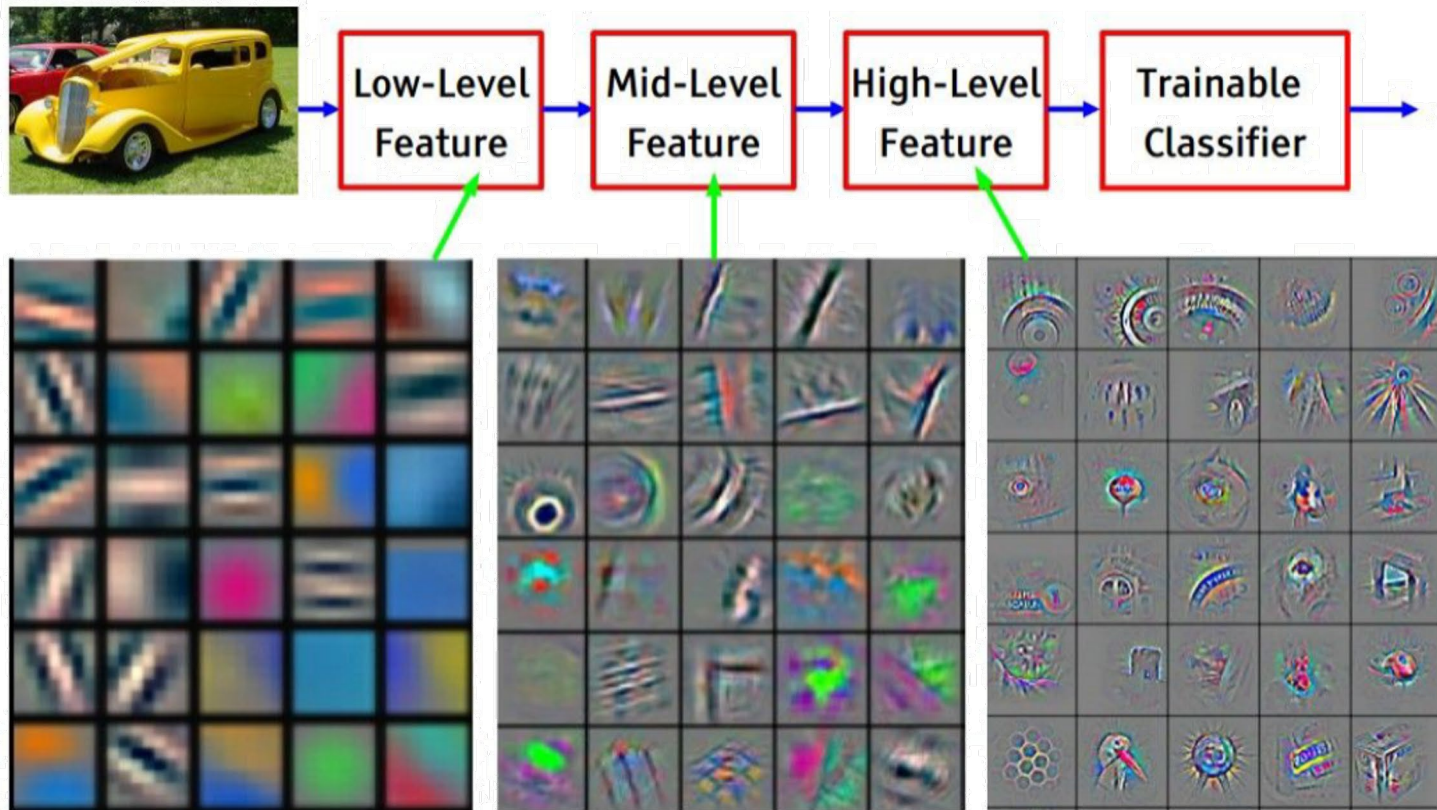
Convolutional Neural Network (CNN)

CNN Prototype

ConvNet is concatenation of Conv Layers and activations




CNN Learned Filters



[Zeiler & Fergus, ECCV'14] Visualizing and Understanding Convolutional Networks

Pooling

Pooling Layer




Pooling Layer: Max Pooling

Single depthslice of input

3	1	3	5
6	0	7	9
3	2	1	4
0	2	4	3

Max pool with
2×2 filters and stride 2




‘Pooled’ output

6	9
3	4

Pooling Layer

- Conv Layer = 'Feature Extraction'
 - Computes a feature in a given region
- Pooling Layer = 'Feature Selection'
 - Picks the strongest activation in a region

Pooling Layer

- Input is a volume of size $W_{in} \times H_{in} \times D_{in}$
- Two hyperparameters
 - Spatial filter extent F
 - Stride S

Filter count K and padding P
make no sense here
- Output volume is of size $W_{out} \times H_{out} \times D_{out}$
 - $W_{out} = \frac{W_{in} - F}{S} + 1$
 - $H_{out} = \frac{H_{in} - F}{S} + 1$
 - $D_{out} = D_{in}$
- Does not contain parameters; e.g. it's fixed function

Pooling Layer

- Input is a volume of size $W_{\text{in}} \times H_{\text{in}} \times D_{\text{in}}$

- Two hyperparameters

- Spatial filter extent F

- Stride S

} Filter count K and padding P
make no sense here

- Output volume is of size $W_{\text{out}} \times H_{\text{out}} \times D_{\text{out}}$

- $W_{\text{out}} = \frac{W_{\text{in}} - F}{S} + 1$

- $H_{\text{out}} = \frac{H_{\text{in}} - F}{S} + 1$

- $D_{\text{out}} = D_{\text{in}}$

- Does not contain parameters; e.g. it's fixed function

Common settings:

$$F = 2, S = 2$$


$$F = 3, S = 2$$

Pooling Layer: Average Pooling

Single depthslice of input

3	1	3	5
6	0	7	9
3	2	1	4
0	2	4	3

Average pool with
2×2 filters and stride 2




‘Pooled’ output

2.5	6
1.75	3

- Typically used deeper in the network

CNN Prototype



Final Fully-Connected Layer

- Same as what we had in ‘ordinary’ neural networks
 - Make the final decision with the extracted features from the convolutions
 - One or two FC layers typically


Convolutions vs Fully-Connected

- In contrast to fully-connected layers, we want to restrict the degrees of freedom
 - FC is somewhat brute force
 - Convolutions are structured
- Sliding window to with the same filter parameters to extract image features
 - Concept of weight sharing
 - Extract same features independent of location


Receptive field

Receptive Field

- Spatial extent of the connectivity of a convolutional filter




5x5 input



3x3 filter


=




3x3 output

Receptive Field

- Spatial extent of the connectivity of a convolutional filter




5x5 input



3x3 filter

=




3x3 output



3x3 receptive field = 1 output pixel is connected to 9 input pixels

Receptive Field

- Spatial extent of the connectivity of a convolutional filter




7x7 input



3x3 receptive field = 1 output pixel is connected to 9 input pixels


Receptive Field

- Spatial extent of the connectivity of a convolutional filter



Receptive Field


- Spatial extent of the connectivity of a convolutional filter





3x3 receptive field = 1 output pixel is connected to 9 input pixels

Receptive Field

- Spatial extent of the connectivity of a convolutional filter



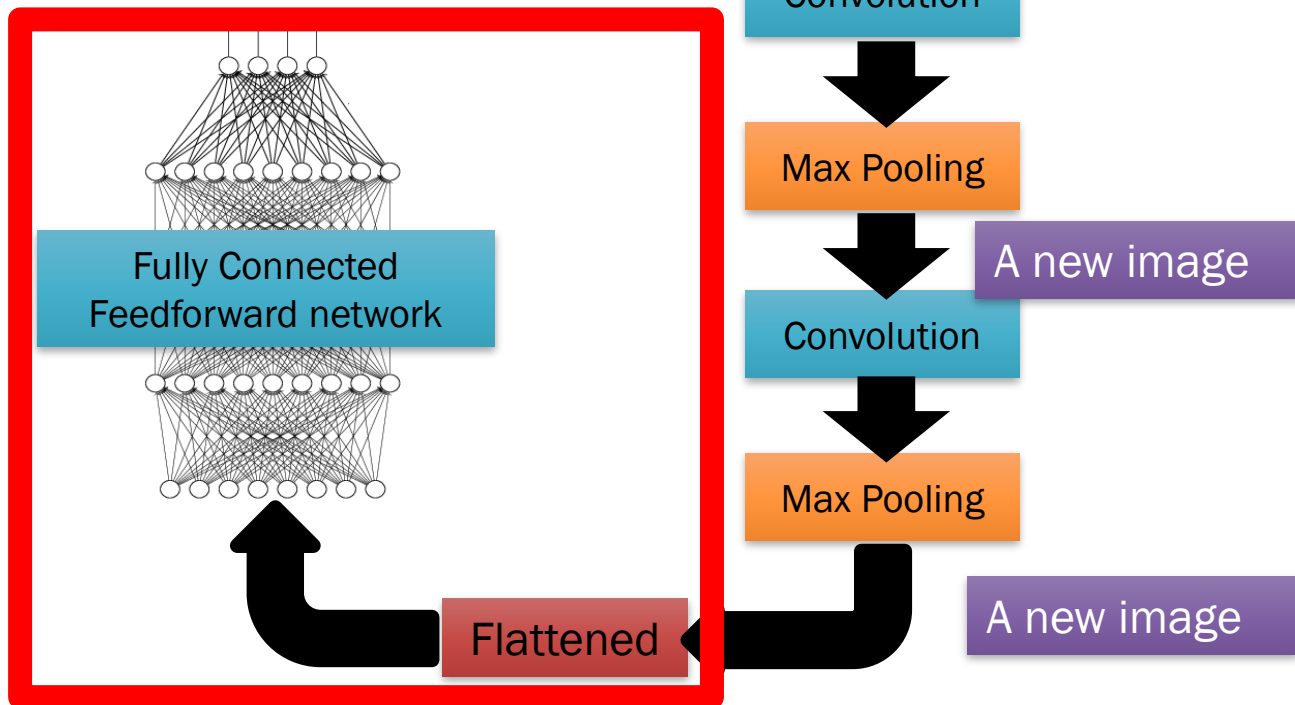
7x7 input



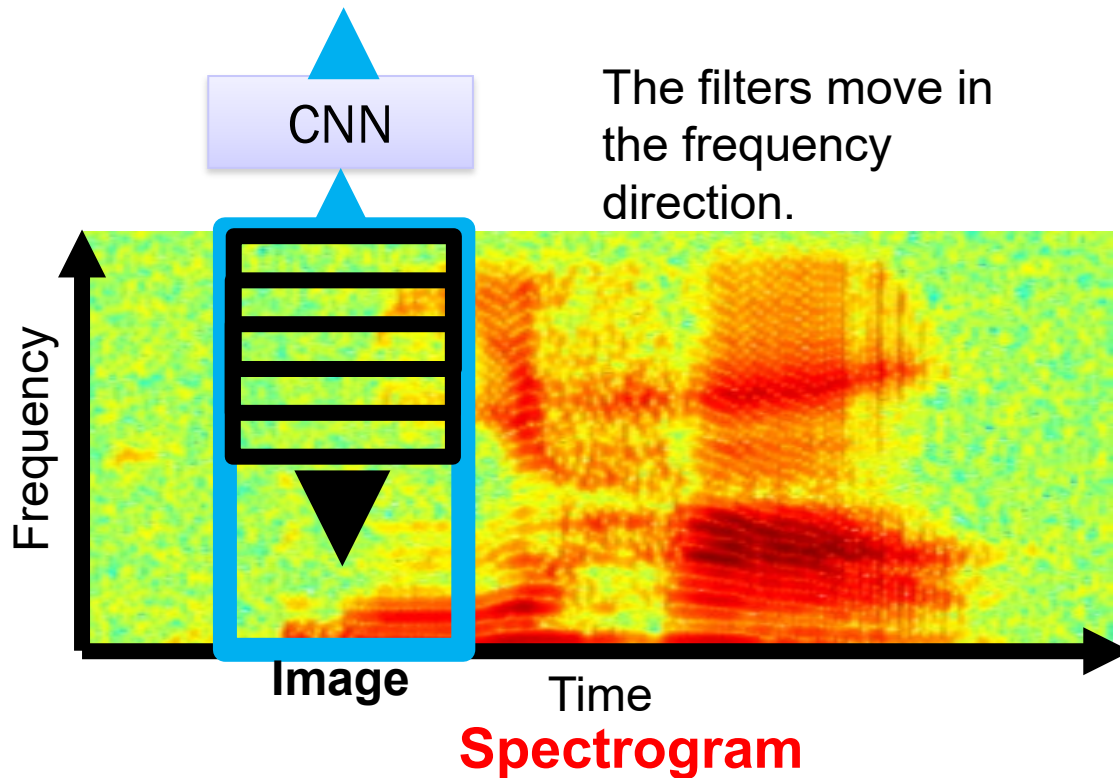
5x5 receptive field on the original input:
one output value is connected to 25 input pixels

The whole CNN

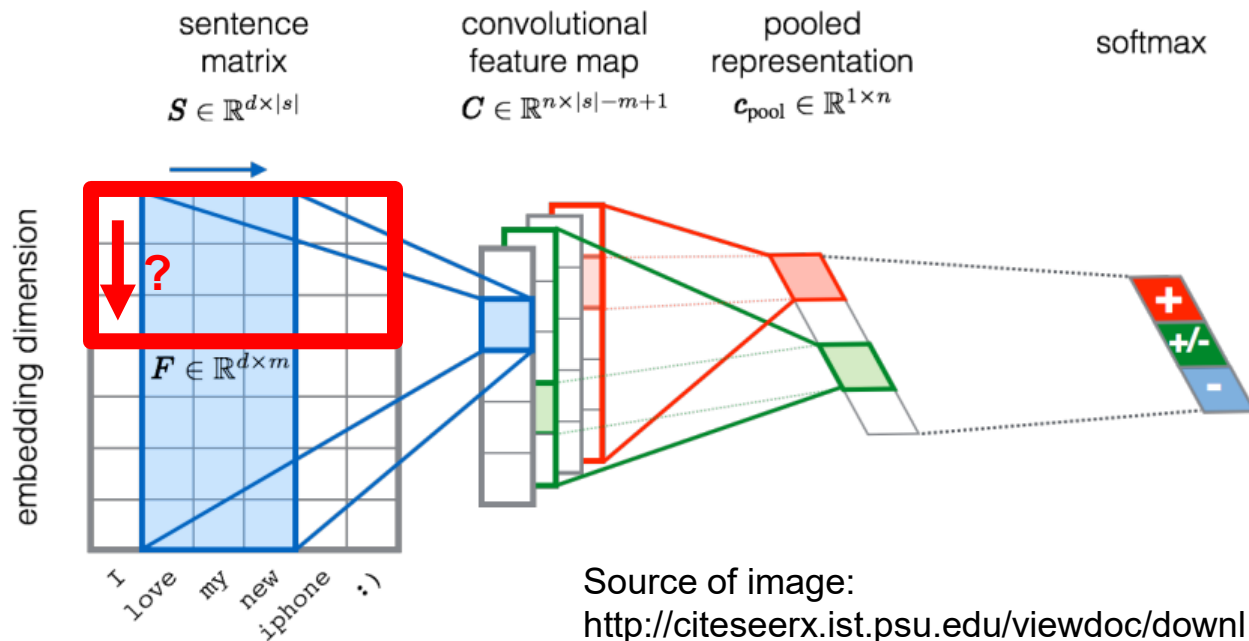
cat dog



CNN in speech recognition

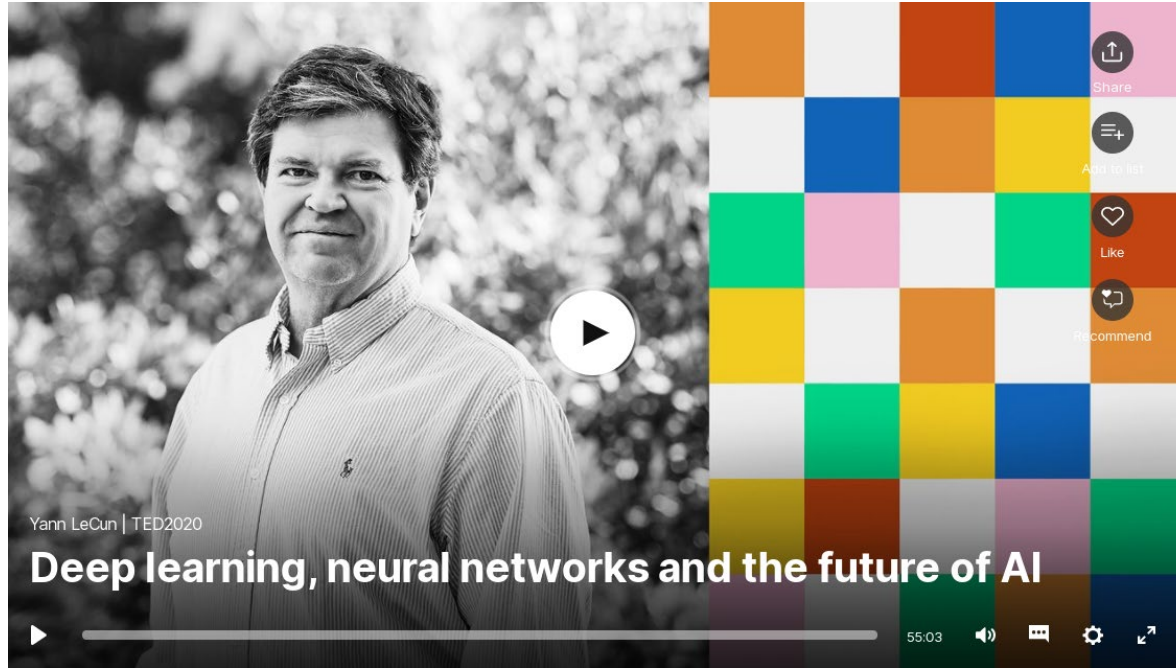


CNN in text classification



Source of image:
<http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.703.6858&rep=rep1&type=pdf>

A lot of buzz about Deep Learning



Yann LeCun, the chief AI scientist at Facebook, helped develop the deep learning algorithms

References

- Goodfellow et al. “Deep Learning” (2016),
 - Chapter 9: Convolutional Networks
- <http://cs231n.github.io/convolutional-networks/>

Acknowledgments

Most slides adapted from [Visual Computing Group](#)